

# CƠ CHẾ AGENT & ENVIRONMENT TRONG KIẾN TRÚC DUAL-AGENT RL IDS

Ngày: 27/10/2025

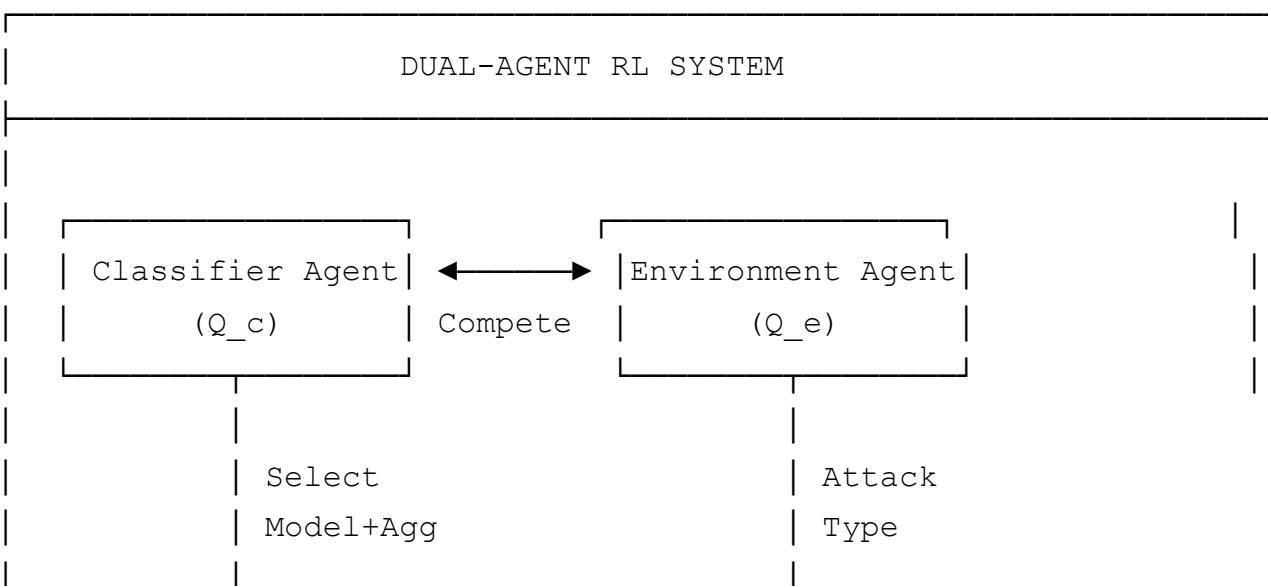
Kiến trúc: Dual-Agent Reinforcement Learning cho Intrusion Detection System

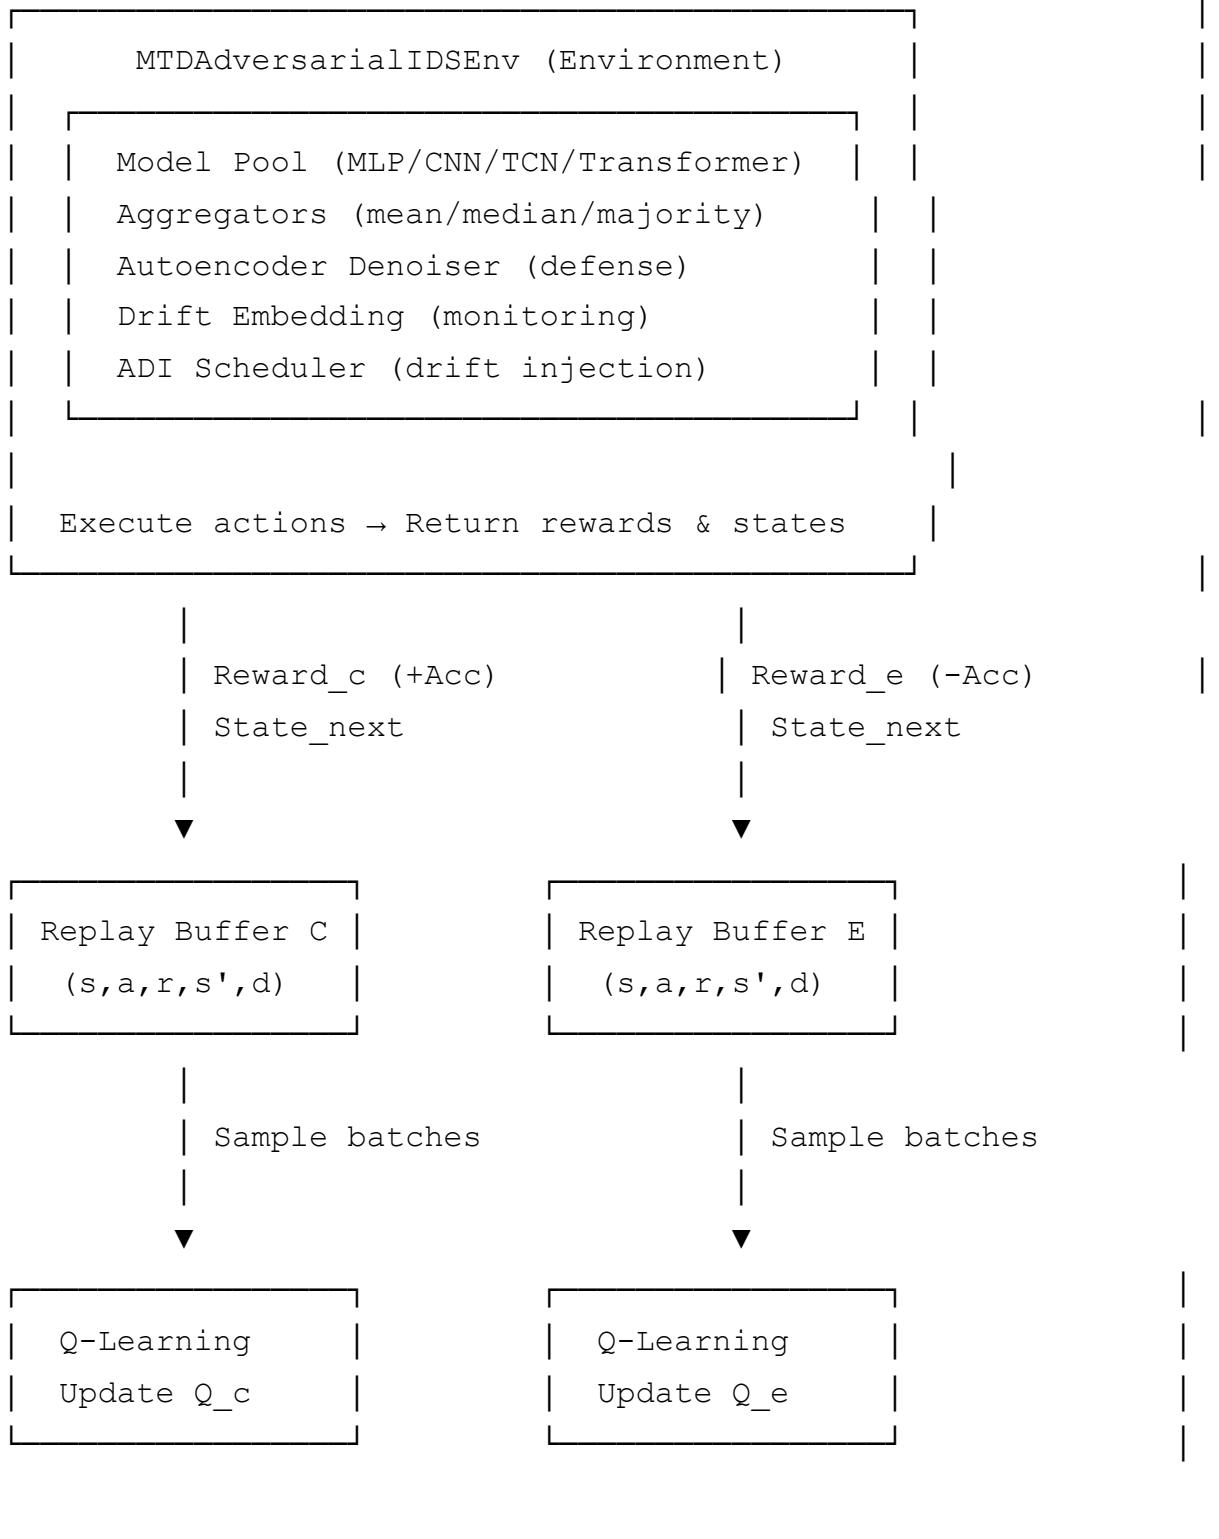
## MỤC LỤC

- [Tổng quan kiến trúc](#)
- [Phân biệt các khái niệm](#)
- [Classifier Agent \( \$Q\_c\$ \)](#)
- [Environment Agent \( \$Q\_e\$ \)](#)
- [MTDAdversarialIDSEnv \(Môi trường\)](#)
- [Luồng tương tác](#)
- [Markov Decision Process \(MDP\)](#)
- [Training Loop](#)
- [So sánh Agent vs Environment](#)

## 1. TỔNG QUAN KIẾN TRÚC

Diagram Tổng Quát:





## 2. PHÂN BIỆT CÁC KHÁI NIỆM

**! QUAN TRỌNG: Có 3 thực thể khác nhau!**

Thực thể	Vai trò	Mục tiêu	Code class
<b>Classifier Agent</b>	RL Agent #1	Maximize accuracy	<b>ClassifierAgent</b>

Thực thể	Vai trò	Mục tiêu	Code class
<b>Environment Agent</b>	RL Agent #2	Minimize accuracy	EnvironmentAgent
<b>Environment (MTD)</b>	Môi trường RL	Execute actions	MTDAdversarialIDSEnv

## 👉 Phân biệt:

### 1. Classifier Agent vs Environment Agent:

- Đây là 2 RL agents **đối nghịch** (adversarial)
- Cùng observe state từ environment
- Có action spaces **RIÊNG BIỆT**
- Có rewards **ĐỐI NGHỊCH**

### 2. Agent vs Environment:

- Agent **quyết định** action
- Environment **thực thi** action và **trả về** (reward, next\_state)
- Relationship: **Actor-Observer** pattern

## 3. CLASSIFIER AGENT (Q\_c)

### 3.1. Định nghĩa:

```
class ClassifierAgent(nn.Module):
    """
    Q_c network: selects (anchor_model, aggregator) pair.
    Action space = |models| × |aggregators|
    """

    def __init__(self, state_dim: int, n_models: int, n_aggregators: int,
                 super().__init__()
                 self.n_models = n_models
                 self.n_aggregators = n_aggregators
                 self.n_actions = n_models * n_aggregators # 4 models × 3 aggs =
# Q-network: State → Q-values for each action
                 self.net = nn.Sequential(
                     nn.Linear(state_dim, hidden),
                     nn.ReLU(),
                     nn.Linear(hidden, hidden),
                     nn.ReLU(),
```

```

        nn.Linear(hidden, self.n_actions), # Output: Q-value for each
    )

```

## 3.2. Action Space:

```

# Example: n_models=4, n_aggregators=3 → 12 actions total
# Action encoding: a_c = anchor_idx + agg_idx * n_models

```

Actions:

0: (MLP, mean)	1: (CNN, mean)	2: (TCN, mean)	3: (Transformer, mean)
4: (MLP, median)	5: (CNN, median)	6: (TCN, median)	7: (Transformer, median)
8: (MLP, majority)	9: (CNN, majority)	10: (TCN, majority)	11: (Transformer, majority)

## 3.3. Decoding Action:

```

def decode_action(self, a_c: int) -> Tuple[int, int]:
    """Decode flat action into (anchor_idx, agg_idx)."""
    # Example: a_c = 7
    # anchor_idx = 7 % 4 = 3 (Transformer)
    # agg_idx = 7 // 4 = 1 (median)
    mid = a_c % self.n_models
    aid = a_c // self.n_models
    return mid, aid

```

## 3.4. Reward Function:

```

# GOAL: MAXIMIZE accuracy
reward_c = w1 · ΔAcc - w2 · Cost - w3 · Switch - w4 · Drift

```

Components:

- ΔAcc: Change in accuracy (higher is better)
- Cost: Computational cost of selected models
- Switch: Penalty for changing models frequently
- Drift: Penalty for high drift magnitude

## 3.5. State Input:

State dimension = 32:

- Recent accuracy history: [10]
- Recent reward history: [10]
- Last actions (normalized): [3] → [model\_idx/n\_models, agg\_idx/n\_aggs,
- Running average reward: [1]
- Drift embedding vector: [11] → KS/PSI/MMD/CUSUM statistics

## 3.6. Learning Process:

# Q-Learning Update:

$$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \cdot \max_{a'} Q(s', a') - Q(s, a)]$$

Where:

- s: current state
  - a: action taken (model+aggregator selection)
  - r: shaped reward (accuracy-based)
  - s': next state
  - $\gamma$ : discount factor (0.95)
- 

## 4. ENVIRONMENT AGENT (Q\_e)

### 4.1. Định nghĩa:

```
class EnvironmentAgent(nn.Module):
    """
    Q_e network: selects adversarial attack (class × eps × attack_type).
    For simplicity, we use a flat action space.
    """

    def __init__(self, state_dim: int, n_attack_actions: int = 12, hidden
                 super().__init__()
                 self.n_actions = n_attack_actions

                 self.net = nn.Sequential(
                     nn.Linear(state_dim, hidden),
                     nn.ReLU(),
                     nn.Linear(hidden, hidden),
                     nn.ReLU(),
                     nn.Linear(hidden, self.n_actions), # 12 attack types
                 )
```

## 4.2. Action Space:

```
# 12 attack actions:  
# Action encoding: a_e = attack_type * 4 + eps_level
```

Attack types:

- 0: Random noise
- 1: Sign-based (FGSM-like)
- 2: Uniform noise

Epsilon levels (perturbation strength):

- Level 0: eps = 0.01
- Level 1: eps = 0.02
- Level 2: eps = 0.03
- Level 3: eps = 0.04

Total actions:

a_e=0: Random, eps=0.01	a_e=4: Sign, eps=0.01	a_e=8: Uniform
a_e=1: Random, eps=0.02	a_e=5: Sign, eps=0.02	a_e=9: Uniform
a_e=2: Random, eps=0.03	a_e=6: Sign, eps=0.03	a_e=10: Uniform
a_e=3: Random, eps=0.04	a_e=7: Sign, eps=0.04	a_e=11: Uniform

## 4.3. Attack Implementation:

```
# In MTDAdversarialIDSEnv.step():  
if a_e is not None and a_e > 0:  
    eps = 0.01 * (1 + (a_e % 4)) # Extract epsilon level  
    attack_type = a_e // 4 # Extract attack type  
  
    if attack_type == 0:  
        # Random perturbation  
        noise = torch.randn_like(x_t) * eps  
    elif attack_type == 1:  
        # Sign-based (FGSM-like without gradient)  
        noise = torch.sign(torch.randn_like(x_t)) * eps  
    else:  
        # Uniform noise  
        noise = (torch.rand_like(x_t) * 2 - 1) * eps  
  
    x_attacked = x_t + noise
```

## 4.4. Reward Function:

```
# GOAL: MINIMIZE accuracy (maximize confusion)
reward_e = -accuracy_raw
```

Where:

- accuracy\_raw = 1.0 if prediction correct, 0.0 if wrong
- Environment Agent wants to create adversarial examples that fool the

## 4.5. State Input:

```
# SAME state as Classifier Agent!
# Both agents observe the same environment state

State dimension = 32 (identical to Q_c)
```

## 4.6. Learning Process:

```
# Q-Learning Update (same algorithm, different reward):
Q_e(s, a) ← Q_e(s, a) + α[r_e + γ · max_a' Q_e(s', a') - Q_e(s, a)]
```

Where:

- r\_e = -accuracy\_raw (negative reward for correct predictions)
- Environment Agent learns to maximize misclassification

---

# 5. MTDAdversarialIDSEnv (MÔI TRƯỜNG)

## 5.1. Định nghĩa:

```
class MTDAdversarialIDSEnv:
    """
    Environment with:
    - Model pool (ensemble)
    - Autoencoder denoiser for adversarial robustness
    - Adversarial perturbation capability
```

- Drift-aware state augmentation

"""

## 5.2. Components:

MTDAdversarialIDSEnv

### 1. MODEL POOL (Ensemble):

- |— MLPDropout (512→256→128→2)
- |— DeepCNN (Conv1D layers)
- |— DeepTCN (Temporal Conv)
- |— Transformer (Self-attention)

### 2. AGGREGATORS:

- |— mean\_logits (Average logits)
- |— median\_logits (Median logits)
- |— majority\_vote (Vote counting)

### 3. AUTOENCODER DENOISER:

- |— Encoder: input → bottleneck (32-dim)
- |— Decoder: bottleneck → reconstructed input
- |— Purpose: Remove adversarial perturbations

### 4. DRIFT EMBEDDING:

- |— KS/PSI statistics (per-feature drift)
- |— MMD (Maximum Mean Discrepancy)
- |— CUSUM (Cumulative Sum detector)

### 5. ADI SCHEDULER:

- |— Covariate drift (feature shift)
- |— Prior drift (label flip)
- |— Conditional drift (region-based flip)

### 6. REWARD SHAPER:

$$\text{r}' = w_1 \cdot \Delta \text{Acc} - w_2 \cdot \text{Cost} - w_3 \cdot \text{Switch} - w_4 \cdot \text{Drift}$$

## 5.3. Key Methods:

### a) reset():

```
def reset(self, seed: Optional[int] = None) -> np.ndarray:  
    """Reset environment to initial state."""  
    self.i = 0 # Reset stream pointer  
    self.hist = [] # Clear history  
    self.adversarial_buffer = [] # Clear attack buffer  
  
    # Initialize state  
    return self._get_state() # Return initial state (32-dim vector)
```

### b) step():

```
def step(self, a_c: int, a_e: Optional[int] = None) -> Tuple[np.ndarray,  
    """  
    Execute one step in environment.  
  
    Args:  
        a_c: Classifier Agent action (model+aggregator selection)  
        a_e: Environment Agent action (attack type, optional)  
  
    Returns:  
        next_state: Next state observation (32-dim)  
        reward_c: Reward for Classifier Agent (shaped)  
        reward_e: Reward for Environment Agent (-accuracy)  
        done: Episode termination flag  
        info: Additional information dict  
    """  
    # 1. Get current data sample  
    x_current = self.X[self.i]  
    y_current = self.y[self.i]  
  
    # 2. Apply ADI drift injection (if active)  
    x_current, y_current, adi_info = self.adi_scheduler.apply(self.i, x_c  
  
    # 3. Apply adversarial attack (if Environment Agent acts)  
    if a_e is not None and a_e > 0:  
        x_attacked = self._apply_attack(x_current, a_e)  
    else:  
        x_attacked = x_current
```

```

# 4. Apply denoiser defense
if self.denoiser is not None:
    x_t = self.denoiser(x_attacked)
else:
    x_t = x_attacked

# 5. Decode Classifier Agent action
anchor_idx = a_c % n_models
agg_idx = a_c // n_models

# 6. Select ensemble subset
selected_models = self._select_ensemble_subset(anchor_idx, model_name)

# 7. Make prediction with ensemble
pred, probs = self._ensemble_predict(x_t, selected_models, agg_method)

# 8. Compute rewards
current_acc = float(pred == y_current)
delta_acc = current_acc - self.prev_accuracy

# Shaped reward for Classifier Agent
reward_c_shaped = self.reward_shaper.total_reward(delta_acc, selected_models)

# Raw accuracy for Environment Agent (negative)
reward_e = -current_acc

# 9. Update drift embedding
self.drift_embedding.push(x_current, score=probs.max(), y=y_current)

# 10. Move to next sample
self.i += 1
done = (self.i >= len(self.X))

# 11. Get next state
next_state = self._get_state() if not done else np.zeros(self.obs_dim)

return next_state, reward_c_shaped, reward_e, done, info

```

### c)\_get\_state():

```

def _get_state(self) -> np.ndarray:
    """Construct state vector (32-dim)."""

```

```

state = np.concatenate([
    np.array(list(self.acc_history), dtype=np.float32),      # [10] rec
    np.array(list(self.reward_history), dtype=np.float32),   # [10] rec
    self.last_actions,                                       # [3] last
    np.array([self.avg_reward], dtype=np.float32),           # [1] runni
    self.drift_embedding.vector() if self.drift_embedding else np.zer
])
return state

```

## 5.4. Action Decoding (CRITICAL):

```

# MUST match ClassifierAgent.decode_action() !

# In Environment:
anchor_idx = a_c_clamped % n_models    # Same as ClassifierAgent
agg_idx = a_c_clamped // n_models      # Same as ClassifierAgent

# CONSISTENCY CHECK:
agent_mid, agent_aid = q_c.decode_action(a_c)
env_mid = a_c % n_models
env_aid = a_c // n_models

assert agent_mid == env_mid  # MUST be equal!
assert agent_aid == env_aid  # MUST be equal!

```

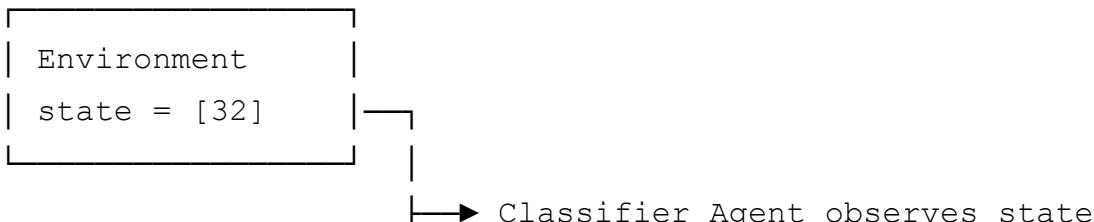
---

# 6. LUÔNG TƯƠNG TÁC

## 6.1. Single Step Flow:

ONE TRAINING STEP

### 1. OBSERVE STATE:



Environment Agent observes state

## 2. SELECT ACTIONS:

```
| Classifier Agent |  
| ε-greedy | → a_c = 7 (Transformer + median)
```

```
| Environment Agent |  
| ε-greedy | → a_e = 5 (Sign attack, eps=0.02)
```

## 3. EXECUTE IN ENVIRONMENT:

```
| MTDAdversarialIDSEnv.step(a_c=7, a_e=5)|  
  
| 1. Get data: x, y |  
| 2. Apply drift: x' = ADI(x) |  
| 3. Apply attack: x'' = Attack(x', a_e) |  
| 4. Denoise: x''' = Denoiser(x'') |  
| 5. Predict: pred = Ensemble(x''', a_c) |  
| 6. Compute rewards:  
|   - reward_c = shaped_reward(...) |  
|   - reward_e = -accuracy |  
| 7. Get next state: s' |
```



```
| RETURNS: (s', r_c, r_e, done, info)
```

## 4. STORE TRANSITIONS:

```
| Replay Buffer C | → (s, a_c, r_c, s', done)
```

```
| Replay Buffer E | → (s, a_e, r_e, s', done)
```

5. LEARN FROM EXPERIENCE:

```
| Sample batch from Replay Buffer C  
| Update Q_c network (Q-learning)
```

```
| Sample batch from Replay Buffer E  
| Update Q_e network (Q-learning)
```

6. REPEAT for next sample...

## 6.2. Episode Flow:

Episode = Processing entire dataset sequentially

```
| EPISODE START (i=0)
```

```
| 1. s = env.reset()
```

▼

```
| STEP LOOP (i=0 to N-1)
```

```
| while not done:
```

```
|   1. Select actions:
```

```
|     a_c = ε-greedy(Q_c(s))
```

```
|     a_e = ε-greedy(Q_e(s))
```

```
|   2. Execute:
```

```
|     s', r_c, r_e, done, info = env.step(a_c, a_e)
```

```
|   3. Store:
```

```
|     replay_c.push(s, a_c, r_c, s', done)
```

```
|     replay_e.push(s, a_e, r_e, s', done)
```

```

| 4. Learn (every few steps):
|   if len(replay_c) >= batch_size:
|     update_network(Q_c, replay_c, ...)
|     update_network(Q_e, replay_e, ...)
|
| 5. Update state:
|   s = s'
|   i += 1

```

| done = True (reached end of dataset)



```

| EPISODE END
| - Decay ε (exploration rate)
| - Save checkpoint
| - Log metrics

```

## 7. MARKOV DECISION PROCESS (MDP)

### 7.1. MDP Framework cho Dual-Agent System:

Hệ thống dual-agent RL của chúng ta có thể được mô hình hóa như một **Multi-Agent Markov Decision Process (MAMDP)** với các thành phần:

$$\text{MDP} = (S, A_c, A_e, P, R_c, R_e, \gamma)$$

Where:

- S : State space (32-dimensional continuous space)
- A\_c : Classifier Agent action space (12 discrete actions)
- A\_e : Environment Agent action space (12 discrete actions)
- P : State transition probability  $P(s'|s, a_c, a_e)$
- R\_c : Reward function for Classifier Agent
- R\_e : Reward function for Environment Agent
- γ : Discount factor (0.95)

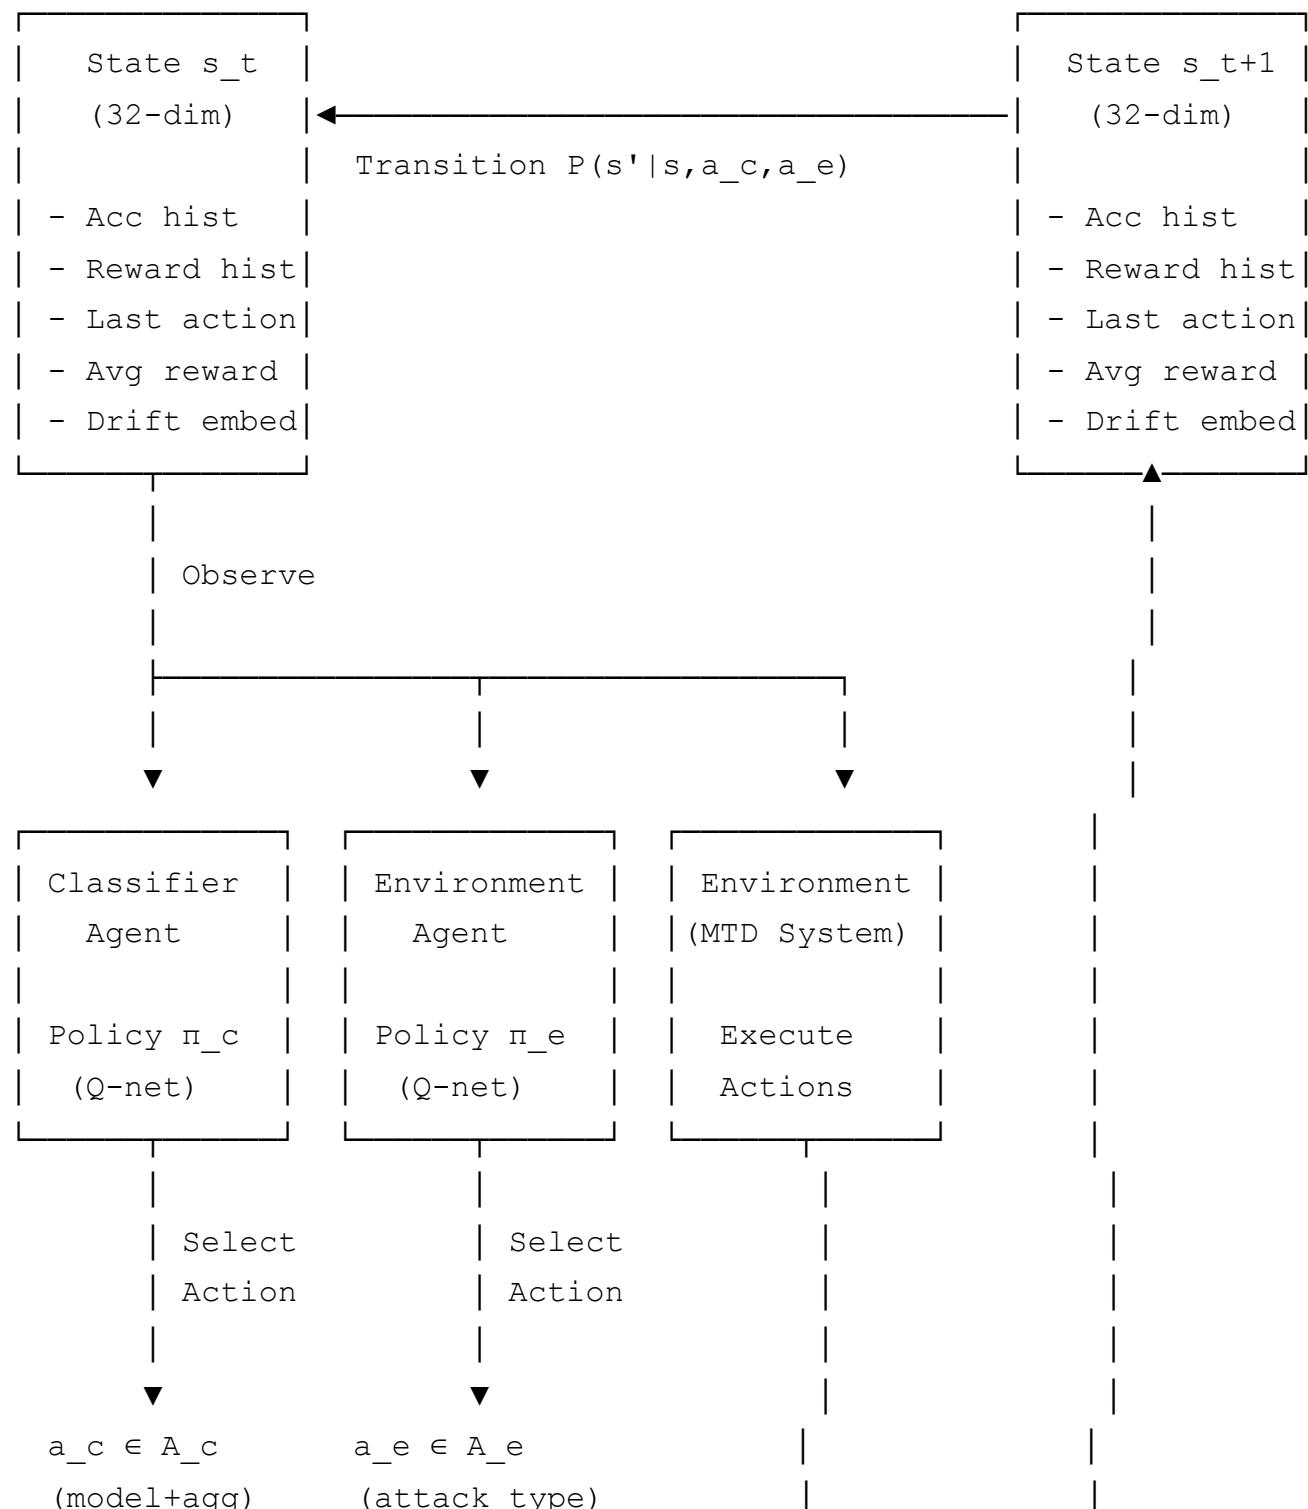
### 7.2. Sơ Đồ MDP Tổng Quan:

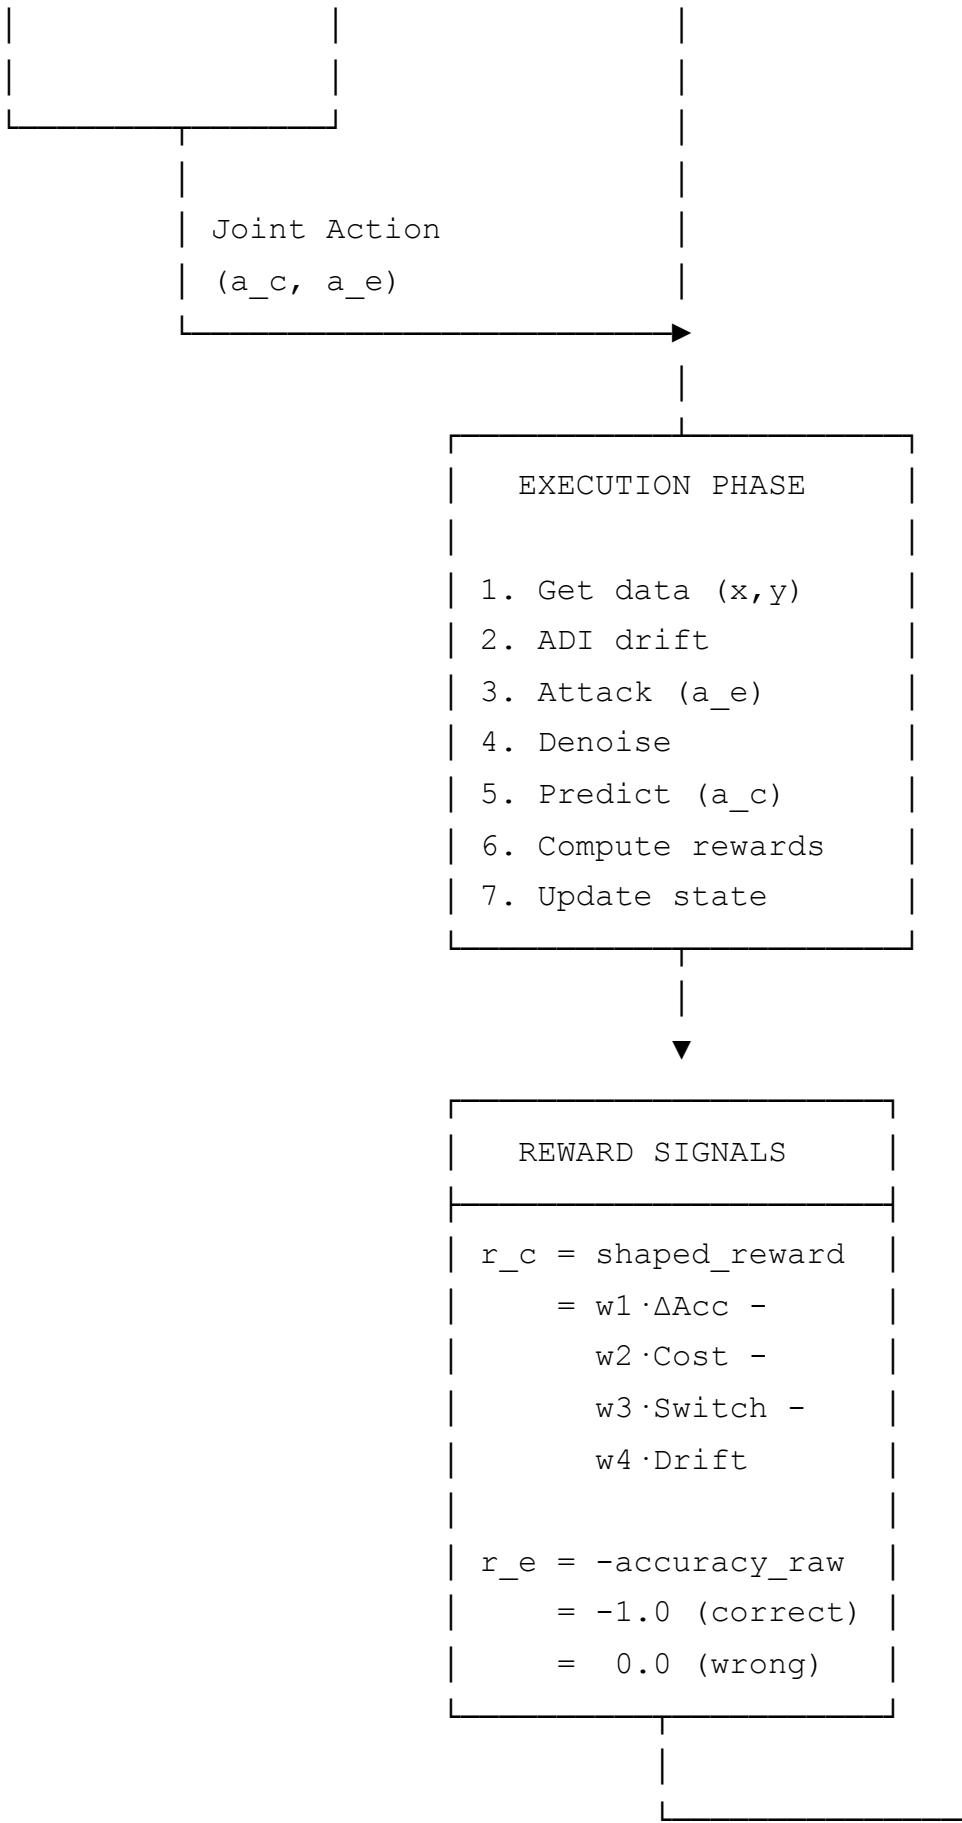
DUAL-AGENT MARKOV DECISION PROCESS

TIME STEP: t

TIME STEP: t+1

2






---



---



---

=

LEARNING UPDATES (Off-policy Q-learning) :

Classifier Agent:

```
Q_c(s_t, a_c) ← Q_c(s_t, a_c) + α[r_c + γ·max_a' Q_c(s_t+1, a') -  
Q_c(s_t, a_c)]
```

Environment Agent:

```
Q_e(s_t, a_e) ← Q_e(s_t, a_e) + α[r_e + γ·max_a' Q_e(s_t+1, a') -  
Q_e(s_t, a_e)]
```

## 7.3. State Space (S):

STATE SPACE S (32-dim)

Dimensions [0-9]: Recent Accuracy History

[acc\_t-9, acc\_t-8, ..., acc\_t-1]

Range: [0.0, 1.0] per dimension

Purpose: Track performance trend

Dimensions [10-19]: Recent Reward History

[r\_t-9, r\_t-8, ..., r\_t-1]

Range: [-1.0, +1.0] (typically)

Purpose: Track reward trajectory

Dimensions [20-22]: Last Actions (Normalized)

[model\_idx/n\_models, agg\_idx/n\_aggs, attack\_idx/12]

Range: [0.0, 1.0] per dimension

Purpose: Action history for temporal patterns

Dimension [23]: Running Average Reward

avg\_reward = EMA(rewards)

Range: [-1.0, +1.0]

Purpose: Long-term performance indicator

Dimensions [24-31]: Drift Embedding (8-dim actual)

- KS/PSI statistics (per-feature drift)
  - MMD (Maximum Mean Discrepancy)
  - CUSUM (Cumulative Sum control chart)
- Range: Normalized [0.0, 1.0]
- Purpose: Detect and quantify concept drift

State Evolution:

$s_t \rightarrow s_{t+1}$  through environment execution

Update rules:

- acc\_history: shift left, append new accuracy
- reward\_history: shift left, append new reward
- last\_actions: update with  $(a_c, a_e)$  normalized
- avg\_reward: exponential moving average
- drift\_embed: update with latest drift statistics

## 7.4. Action Spaces ( $A_c, A_e$ ):

CLASSIFIER AGENT ACTION SPACE ( $A_c$ )

$|A_c| = 12$  discrete actions

Action encoding:  $a_c = \text{anchor\_idx} + \text{agg\_idx} \times n_{\text{models}}$

Anchor Model	Aggregator	Action ID	Effect
MLP (0)	Mean (0)	0	Fast
CNN (1)	Mean (0)	1	Balanced
TCN (2)	Mean (0)	2	Temporal
Transformer (3)	Mean (0)	3	Accurate
MLP (0)	Median (1)	4	Robust

CNN (1)	Median (1)	5	Robust	
TCN (2)	Median (1)	6	Robust	
Transformer (3)	Median (1)	7	Robust	
MLP (0)	Majority(2)	8	Vote-based	
CNN (1)	Majority(2)	9	Vote-based	
TCN (2)	Majority(2)	10	Vote-based	
Transformer (3)	Majority(2)	11	Vote-based	

Policy:  $\pi_c(a_c|s) = \epsilon\text{-greedy over } Q_c(s, \cdot)$

#### ENVIRONMENT AGENT ACTION SPACE ( $A_e$ )

$|A_e| = 12$  discrete actions

Action encoding:  $a_e = \text{attack\_type} \times 4 + \text{eps\_level}$

Attack Type	Epsilon	Action ID	Strength
Random (0)	0.01	0	Very Weak
Random (0)	0.02	1	Weak
Random (0)	0.03	2	Medium
Random (0)	0.04	3	Strong
Sign-based (1)	0.01	4	Weak FGSM
Sign-based (1)	0.02	5	Medium FGSM
Sign-based (1)	0.03	6	Strong FGSM
Sign-based (1)	0.04	7	Very Strong
Uniform (2)	0.01	8	Weak Uniform
Uniform (2)	0.02	9	Medium Uniform
Uniform (2)	0.03	10	Strong Uniform
Uniform (2)	0.04	11	Max Uniform

Policy:  $\pi_e(a_e|s) = \epsilon\text{-greedy over } Q_e(s, \cdot)$

## 7.5. State Transition Function $P(s'|s, a_c, a_e)$ :

### STATE TRANSITION DYNAMICS

```
P(s_{t+1} | s_t, a_c, a_e) = Pr[next state | current context]
```

Transition Process:

#### 1. DETERMINISTIC COMPONENTS:

- Data stream:  $x_{t+1}, y_{t+1} \leftarrow \text{Dataset}[i+1]$
- History update: shift and append
- Action recording:  $\text{last\_actions} \leftarrow (a_c, a_e)$

#### 2. STOCHASTIC COMPONENTS:

- ADI drift:  $x' \sim \text{DriftInjection}(x, \text{schedule})$
- Attack noise:  $x'' \sim \text{Attack}(x', a_e, \varepsilon)$
- Model prediction:  $\text{pred} \sim \text{Ensemble}(x'', a_c)$
- Accuracy:  $\text{acc} = [\text{pred} == y]$  (stochastic via model)

#### 3. STATE CONSTRUCTION:

```
s_{t+1} = [
    acc_history[1:] + [acc_new],      # Shift & append
    reward_history[1:] + [r_new],     # Shift & append
    normalize(a_c, a_e),            # Action encoding
    EMA(rewards),                  # Running average
    DriftEmbedding.vector()         # Drift stats
]
```

Markov Property:

```
P(s_{t+1}|s_t, a_c, a_e) = P(s_{t+1}|s_t, a_c, a_e, s_{t-1}, ...)
```

→ State  $s_t$  contains sufficient statistics (history buffer)

## 7.6. Reward Functions ( $R_c, R_e$ ):

## REWARD FUNCTIONS

$R_c: S \times A_c \times A_e \rightarrow \mathbb{R}$  (Classifier Agent reward)

---

$$r_c(s, a_c, a_e) = w_1 \cdot \Delta \text{Acc} - w_2 \cdot \text{Cost} - w_3 \cdot \text{Switch} - w_4 \cdot \text{Drift}$$

Component	Formula	Weight	Range
Accuracy $\Delta$	$acc_t - acc_{t-1}$	$w_1=1.0$	$[-1, +1]$
Cost Penalty	$\sum \text{model\_complexity} / \max$	$w_2=0.2$	$[0, 1]$
Switch Cost	[action changed]	$w_3=0.1$	$\{0, 1\}$
Drift Impact	$\ \text{drift\_vector}\ $	$w_4=0.3$	$[0, 1]$

Design Goal: Maximize long-term cumulative return

$$G_t^c = \sum_{k=0}^{\infty} \gamma^k \cdot r_c(s_{t+k}, a_c, a_e)$$

Optimal Policy:

$$\pi^*_c = \operatorname{argmax}_{\{\pi_c\}} [G_t^c | \pi_c, \pi_e]$$

$R_e: S \times A_c \times A_e \rightarrow \mathbb{R}$  (Environment Agent reward)

---

$$r_e(s, a_c, a_e) = -\text{accuracy\_raw}$$

Outcome	Accuracy	Reward $r_e$	Interpretation
Correct predict	1.0	-1.0	BAD (defense)
Wrong predict	0.0	0.0	GOOD (attack)

Design Goal: Minimize classifier accuracy (adversarial)

$$\begin{aligned} G_t^e &= \sum_{k=0}^{\infty} \gamma^k \cdot r_e(s_{t+k}, a_c, a_e) \\ &= -\sum_{k=0}^{\infty} \gamma^k \cdot \text{accuracy}_{t+k} \end{aligned}$$

Optimal Policy:

$$\begin{aligned}\pi^*_e &= \operatorname{argmax}_{\{\pi_e\}} [G_t^e \mid \pi_c, \pi_e] \\ &= \operatorname{argmin}_{\{\pi_e\}} [\Sigma \text{ accuracy} \mid \pi_c, \pi_e]\end{aligned}$$

ZERO-SUM GAME PROPERTY:

While  $r_c + r_e \neq 0$  (not strictly zero-sum due to shaped reward),  
the agents have OPPOSING objectives:

- Classifier: maximize accuracy
- Attacker: minimize accuracy

This creates adversarial pressure for robustness.

## 7.7. Bellman Equations:

### BELLMAN OPTIMALITY EQUATIONS

For Classifier Agent ( $Q_c$ ):

$$Q^*_c(s, a_c) = [r_c(s, a_c, a_e) + \gamma \cdot \max_{\{a'_c\}} Q^*_c(s', a'_c)]$$

Iterative Update (Q-learning):

$$Q_c(s_t, a_c) \leftarrow Q_c(s_t, a_c) + \alpha [r_c + \gamma \cdot \max_{\{a'\}} Q_c(s_{t+1}, a') - Q_c(s_t, a_c)]$$

TD Error  $\delta_c$

Where:

$\alpha$  = learning rate (e.g., 0.001 via Adam optimizer)

$\gamma$  = discount factor (0.95)

$\delta_c$  = temporal difference error

Convergence Guarantee:

If all  $(s, a)$  pairs visited infinitely often and learning rate satisfies Robbins-Monro conditions, then  $Q_c \rightarrow Q^*_c$  as  $t \rightarrow \infty$

For Environment Agent ( $Q_e$ ):

$$Q^*_e(s, a_e) = [r_e(s, a_c, a_e) + \gamma \cdot \max_{\{a'\}_e} Q^*_e(s', a'_e)]$$

Iterative Update (Q-learning):

$$Q_e(s_t, a_e) \leftarrow Q_e(s_t, a_e) + \alpha[r_e + \gamma \cdot \max_{\{a'\}} Q_e(s_{t+1}, a') - Q_e(s_t, a_e)]$$

TD Error  $\delta_e$

Independent Learning:

- Both agents learn simultaneously
- Non-stationary environment (each agent's policy changes)
- No Nash equilibrium guarantee in general
- Empirically: adversarial training improves robustness

## 7.8. Trajectory và Episode:

### TRAJECTORY STRUCTURE

Episode = One complete pass through dataset

Length = N samples in streaming data

Trajectory  $\tau$ :

```
 $\tau = (s_0, a_{c0}, a_{e0}, r_{c0}, r_{e0},$ 
       $s_1, a_{c1}, a_{e1}, r_{c1}, r_{e1},$ 
       $\dots,$ 
       $s_N, a_{cN}, a_{eN}, r_{cN}, r_{eN})$ 
```

Cumulative Returns:

$$G_0^C = \sum_{t=0}^{N-1} \gamma^t \cdot r_{ct} \quad (\text{Classifier})$$

$$G_0^E = \sum_{t=0}^{N-1} \gamma^t \cdot r_{et} \quad (\text{Environment})$$

Example Episode Flow:

```
t=0: s0 → (a_c=3, a_e=5) → r_c=+0.12, r_e=-1.0 → s1
t=1: s1 → (a_c=7, a_e=8) → r_c=+0.05, r_e=-1.0 → s2
t=2: s2 → (a_c=7, a_e=11) → r_c=-0.20, r_e= 0.0 → s3
...
t=N: s_N → TERMINAL STATE
```

Properties:

- Finite horizon:  $T = N$  (dataset size)
- Deterministic termination:  $\text{done} = (i \geq N)$
- No early stopping (process all samples)
- Episodic task (clear start and end)

## 7.9. Nash Equilibrium và Convergence:

GAME-THEORETIC ANALYSIS

Game Type: Two-Player General-Sum Markov Game

Players:

- Player 1: Classifier Agent (maximizes  $r_c$ )

- Player 2: Environment Agent (maximizes  $r_e = -\text{accuracy}$ )

Strategy Space:

- $\Pi_C = \{\pi_C : S \rightarrow \Delta(A_C)\}$  (policies for Classifier)
- $\Pi_E = \{\pi_E : S \rightarrow \Delta(A_E)\}$  (policies for Environment)

Nash Equilibrium ( $\pi^*_C, \pi^*_E$ ):

$$\begin{aligned} V_C(s; \pi^*_C, \pi^*_E) &\geq V_C(s; \pi_C, \pi^*_E) \quad \forall \pi_C, \forall s \\ V_E(s; \pi^*_C, \pi^*_E) &\geq V_E(s; \pi^*_C, \pi_E) \quad \forall \pi_E, \forall s \end{aligned}$$

Interpretation:

Neither agent can improve by unilaterally changing its policy given the other agent's fixed policy.

Convergence Challenges:

Non-stationary Environment:

- Each agent's policy changes during training
- Violates MDP stationarity assumption
- Can lead to oscillations or cycles

No Guaranteed Convergence:

- Independent Q-learning may not converge to Nash
- Can converge to suboptimal equilibria

✓ Empirical Success:

- Adversarial training improves robustness in practice
- Experience replay stabilizes learning
- Slow policy updates reduce non-stationarity

Practical Goal:

Find approximate equilibrium  $(\hat{\pi}_C, \hat{\pi}_E)$  where:

- Classifier achieves high accuracy under attack
- Attacker forces classifier to be robust
- Co-evolution leads to better generalization

# 8. TRAINING LOOP

## 7.1. Main Training Function:

```
def train_dual_agent_rl(
    env: MTDAAdversarialIDSEnv,
    q_c: ClassifierAgent,
    q_e: EnvironmentAgent,
    replay_c: ReplayBuffer,
    replay_e: ReplayBuffer,
    optimizer_c: torch.optim.Optimizer,
    optimizer_e: torch.optim.Optimizer,
    prequential: Optional[PrequentialMetrics] = None,
    start_step: int = 0,
    epochs: int = 3,
    batch_size: int = 64,
    gamma: float = 0.95,
    epsilon: float = 0.2,
) -> Tuple[int, float]:
    """Train dual-agent RL system."""

    device = env.device
    global_step = start_step

    for epoch in range(epochs):
        # Reset environment for new episode
        state = env.reset(seed=epoch)
        done = False

        while not done:
            # 1. SELECT ACTIONS ( $\varepsilon$ -greedy)
            if random.random() < epsilon:
                a_c = random.randint(0, q_c.n_actions - 1)
                a_e = random.randint(0, q_e.n_actions - 1)
            else:
                with torch.no_grad():
                    s_t = torch.from_numpy(state).unsqueeze(0).to(device)
                    a_c = q_c(s_t).argmax().item()
                    a_e = q_e(s_t).argmax().item()
```

```

# 2. EXECUTE in environment
next_state, reward_c, reward_e, done, info = env.step(a_c, a_e)

# 3. STORE transitions
replay_c.push(state, a_c, reward_c, next_state, done)
replay_e.push(state, a_e, reward_e, next_state, done)

# 4. LEARN from experience
if len(replay_c) >= batch_size and global_step % 4 == 0:
    _update_network(q_c, replay_c, optimizer_c, batch_size, c)
    _update_network(q_e, replay_e, optimizer_e, batch_size, c)

# 5. UPDATE state
state = next_state
global_step += 1

# 6. CHECKPOINT & LOGGING
if global_step % SAVE_EVERY_STEPS == 0:
    save_checkpoint(q_c, q_e, optimizer_c, optimizer_e, global_step)

# Decay exploration
epsilon = max(0.01, epsilon * 0.95)

return global_step, epsilon

```

## 7.2. Q-Learning Update:

```

def _update_network(
    q_net: nn.Module,
    replay: ReplayBuffer,
    optimizer: torch.optim.Optimizer,
    batch_size: int,
    gamma: float,
    device: torch.device,
):
    """Update Q-network using experience replay."""

    # 1. SAMPLE batch from replay buffer
    s, a, r, ns, done = replay.sample(batch_size)

    # 2. Convert to tensors

```

```

s_t = torch.from_numpy(s).to(device)
a_t = torch.from_numpy(a).to(device)
r_t = torch.from_numpy(r).to(device)
ns_t = torch.from_numpy(ns).to(device)
done_t = torch.from_numpy(done).to(device)

# 3. COMPUTE current Q-values
q_pred = q_net(s_t).gather(1, a_t.view(-1, 1)).squeeze(1)

# 4. COMPUTE target Q-values (Bellman equation)
with torch.no_grad():
    q_next = q_net(ns_t).max(dim=1)[0] # Max Q-value for next state
    target = r_t + gamma * q_next * (1.0 - done_t) # TD target

# 5. COMPUTE loss
loss = F.smooth_l1_loss(q_pred, target)

# 6. BACKPROPAGATION
optimizer.zero_grad()
loss.backward()
torch.nn.utils.clip_grad_norm_(q_net.parameters(), max_norm=1.0)
optimizer.step()

```

---

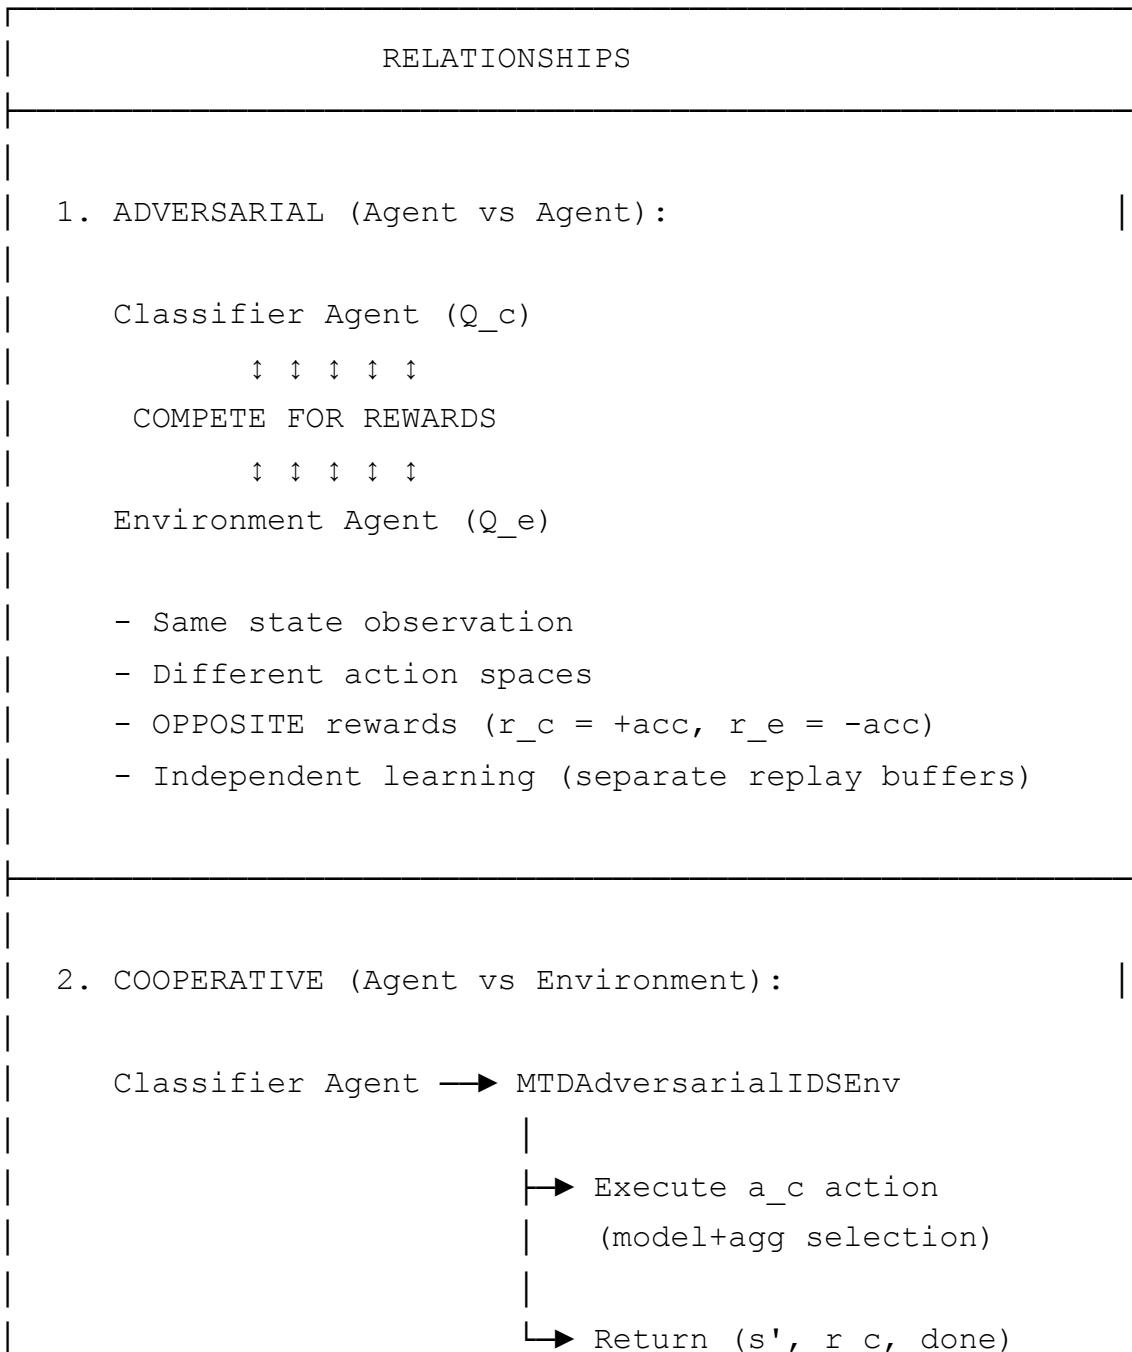
## 9. SO SÁNH AGENT VS ENVIRONMENT

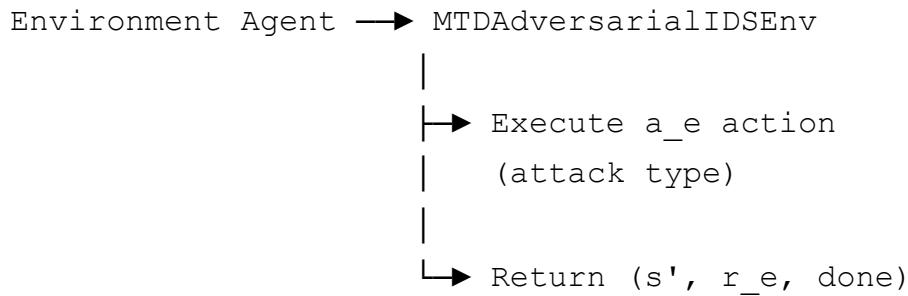
### 8.1. Bảng So Sánh Chi Tiết:

Khía cạnh	Classifier Agent (Q_c)	Environment Agent (Q_e)	MTDAdversarialIDSEnv
Vai trò	RL Agent (Defense)	RL Agent (Attack)	Môi trường RL
Mục tiêu	Maximize accuracy	Minimize accuracy	Execute & return results
Action Space	12 actions (4x3)	12 actions (attack types)	N/A
Action Type	Model+Aggregator selection	Attack type selection	Execute actions
Reward	+Shaped ( $\Delta$ Acc-Cost-Switch-Drift)	-Accuracy (raw)	Compute & return rewards
State	Observe 32-dim vector	Observe 32-dim vector	Provide state vector

Khía cạnh	Classifier Agent ( $Q_c$ )	Environment Agent ( $Q_e$ )	MTDAdversarialIDSEnv
<b>Learning</b>	Q-learning (maximize reward)	Q-learning (maximize - acc)	No learning
<b>Network</b>	MLP (state → Q-values)	MLP (state → Q-values)	Ensemble models
<b>Replay Buffer</b>	Separate (capacity=100k)	Separate (capacity=100k)	N/A
<b>Exploration</b>	$\epsilon$ -greedy	$\epsilon$ -greedy	N/A
<b>Update Frequency</b>	Every 4 steps	Every 4 steps	Every step

## 8.2. Relationship Diagram:

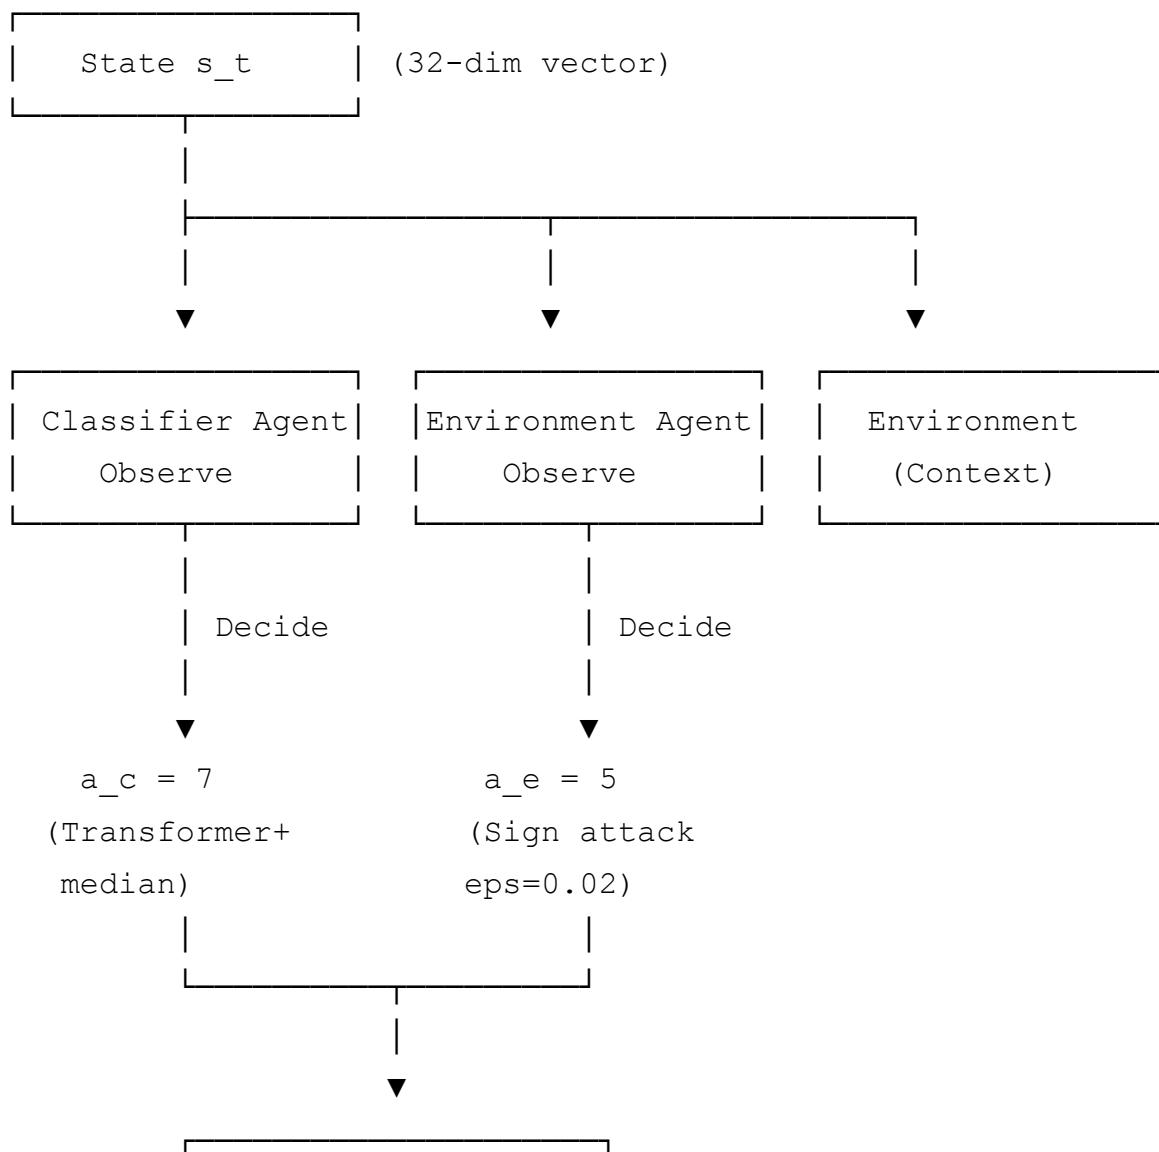


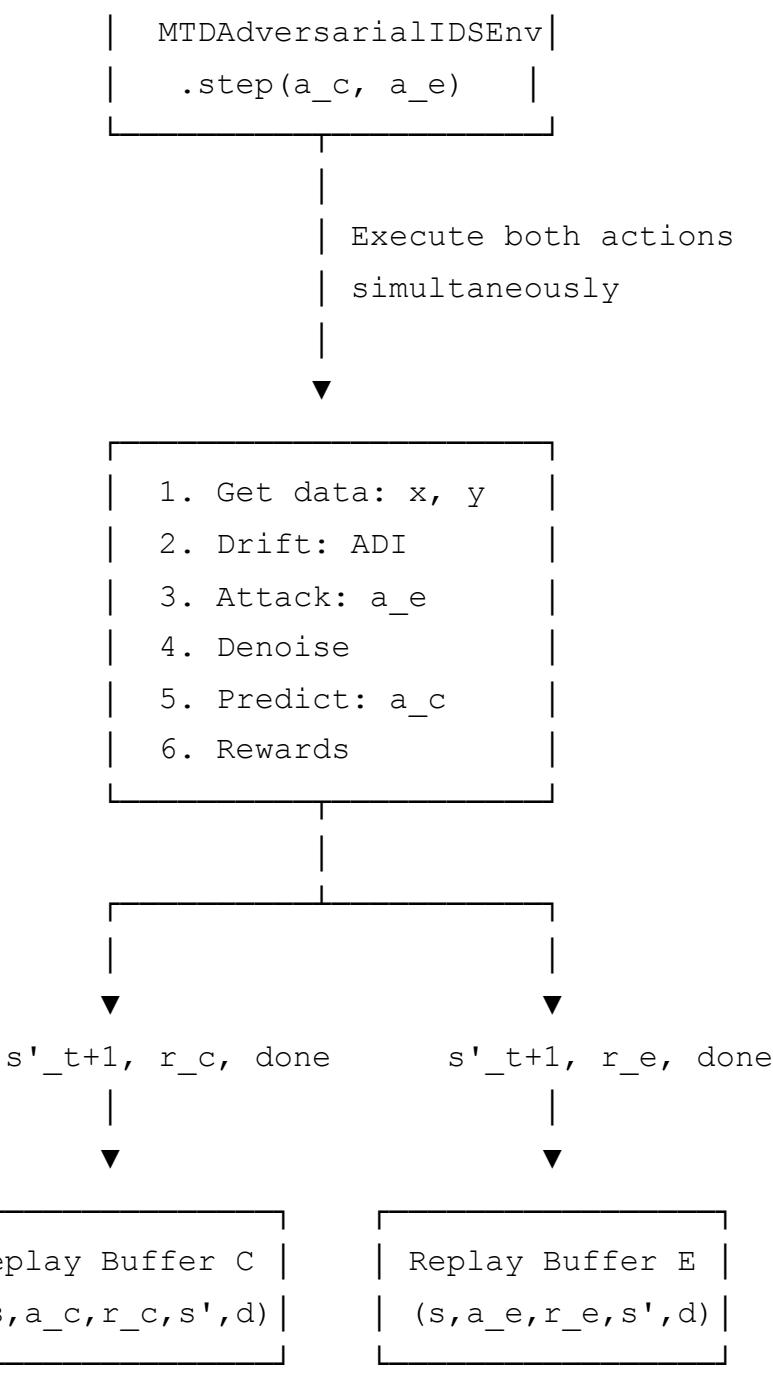


- Environment MUST correctly decode actions
- Consistent action space definition
- Environment provides state observations

### 8.3. Information Flow:

TIME: t





## 9. CÁC ĐIỂM QUAN TRỌNG

### 9.1. Action Space Consistency:

```

# ✅ CORRECT (Current implementation):
# ClassifierAgent.decode_action():
mid = a_c % n_models
aid = a_c // n_models

# MTDAdversarialIDSEnv.step():
anchor_idx = a_c % n_models # SAME formula
agg_idx = a_c // n_models   # SAME formula
    
```

```

# Result: Agent learns correct mapping between actions and outcomes

# ❌ WRONG (Before fix):
# ClassifierAgent.decode_action():
mid = a_c % n_models
aid = a_c // n_models

# MTDAversarialIDSEnv.step():
anchor_idx = a_c % n_models
agg_idx = (a_c // n_models) % n_aggs # DIFFERENT formula (extra modulo)

# Result: Agent learns wrong mapping! Action 11 decoded differently!

```

## 9.2. Reward Signals:

```

# Classifier Agent: Maximize shaped reward
reward_c = w1 · ΔAcc - w2 · Cost - w3 · Switch - w4 · Drift
    = 1.0 · (+0.1) - 0.2 · (0.3) - 0.1 · (0.2) - 0.3 · (0.05)
    = +0.1 - 0.06 - 0.02 - 0.015
    = +0.005 # Small positive reward

# Environment Agent: Maximize negative accuracy
reward_e = -accuracy_raw
    = -1.0 (if correct prediction)
    = -0.0 (if wrong prediction)
# Environment Agent wants prediction to be WRONG!

```

## 9.3. Separate Replay Buffers:

```

# WHY separate buffers?
# 1. Different action spaces (model+agg vs attack type)
# 2. Different rewards (shaped vs raw negative)
# 3. Independent learning (no interference)
# 4. Different exploration strategies (can have different  $\epsilon$ )

replay_c = ReplayBuffer(capacity=100000) # For Classifier Agent
replay_e = ReplayBuffer(capacity=100000) # For Environment Agent

```

## 9.4. Exploration Strategy:

```
# ε-greedy exploration (shared ε for both agents):
if random.random() < epsilon:
    # EXPLORE: Random action
    a_c = random.randint(0, q_c.n_actions - 1)
    a_e = random.randint(0, q_e.n_actions - 1)
else:
    # EXPLOIT: Best known action
    a_c = q_c(state).argmax()
    a_e = q_e(state).argmax()

# Decay schedule:
epsilon = max(0.01, epsilon * 0.95) # Decay by 5% each epoch
```

---

# 10. TÓM TẮT

## 10.1. Key Takeaways:

### 1. 3 thực thể khác nhau:

- Classifier Agent ( $Q_c$ ): Defense player
- Environment Agent ( $Q_e$ ): Attack player
- MTDAdversarialIDSEnv: Game referee

### 2. Adversarial relationship (Agent vs Agent):

- Opposite rewards
- Same state observation
- Different action spaces
- Compete to improve each other

### 3. Cooperative relationship (Agent vs Environment):

- Consistent action decoding
- Environment executes actions faithfully
- Returns rewards and next states
- Provides state observations

### 4. Learning mechanism:

- Q-learning for both agents

- Experience replay for stability
- $\epsilon$ -greedy exploration
- Separate replay buffers

## 5. Architecture goal:

- Robust IDS through adversarial training
- Adaptive model selection (MTD)
- Drift-aware learning
- Efficient ensemble management

## 10.2. Workflow Summary:

### 1. Initialize:

- Create environment with model pool
- Create 2 agents ( $Q_c$ ,  $Q_e$ )
- Create 2 replay buffers

### 2. Training loop:

FOR each epoch:

    Reset environment

    FOR each sample in dataset:

1. Both agents observe state
2. Both agents select actions ( $\epsilon$ -greedy)
3. Environment executes both actions
4. Environment returns rewards & next state
5. Store transitions in separate buffers
6. Update networks from replay buffers
7. Move to next state

    Decay exploration rate

### 3. Evaluation:

- Test on holdout set
- Measure accuracy, F1, AUC
- Test adversarial robustness
- Test drift detection

## 11. REFERENCES

- **Q-Learning:** Watkins & Dayan (1992)
- **Experience Replay:** Lin (1992)

- **Adversarial Training:** Goodfellow et al. (2014)
  - **Moving Target Defense:** Zhuang et al. (2019)
  - **Drift Detection:** Gama et al. (2014)
- 

 **Ghi chú:** Document này mô tả kiến trúc hiện tại AFTER fixes. Trước khi fix, action decoding không consistent giữa Agent và Environment, gây lỗi học sai.